



dpplr Interfaces to Large-Scale Data

Ian Cook

@ianmcook

ian@cloudera.com

Context


Mission for Cloudera: Provide a platform for data analysts, data scientists to efficiently query, analyze, model large-scale data in clusters, cloud storage

- By distributing Apache Spark, Apache Impala, other tools
- By enabling productive use of these tools


Python and R users often have difficulty moving from smaller data to large-scale distributed data

- Familiar packages, methods don't work the same way on distributed data

Poll question



Poll question




dplyr

dplyr provides a set of *verbs* that perform common data manipulation steps

- `select()` to select columns
- `filter()` to filter rows
- `arrange()` to order rows
- `mutate()` to create new columns
- `summarise()` to aggregate
- `group_by()` to perform operations by group

dplyr works on local data and with remote data sources

- For remote sources, dplyr commands are translated into SQL



Poll question

Demonstration

Example code at
github.com/ianmcook/dplyr-examples

dplyr SQL backends



* optional

sparklyr

- Provides a SQL backend to dplyr for Spark
- Also exposes the MLlib API and a subset of the Spark DataFrames API
- Developed by RStudio


spark.rstudio.com



implyr

- Provides a SQL backend to dplyr for Impala
- Uses ODBC or JDBC to connect to Impala
- Developed at Cloudera

tiny.cloudera.com/implyr



Five tips for using dplyr with SQL data sources

1

Use `show_query()`

2

`filter()` early
`arrange()` late

3

Check your data types

4

Know your SQL engine

5

Know when to collect ()

Cloudera Data Science Workbench

More information

tiny.cloudera.com/cdsw

OnDemand training

tiny.cloudera.com/cdsw-training

